

COMS 4771: Machine Learning (Fall 2018) - Homework #3

Amal Alabdulkarim (aa4235), Danyang He (dh2914), Jing Qian (jq2282)

November 20, 2018

Problem 1

(i)

Given (y, X) and (α, λ) , we can augment the data by adding d more “fake data points”: $\sqrt{\lambda\alpha}\vec{e}_i$ ($\vec{e}_i \in R^d$, represents the vector having only i^{th} element nonzero) and augment the label by adding d more 0 labels. Thus, the augmented data matrix $X^* \in R^{d \times (n+d)}$, $y^* \in R^{1 \times (n+d)}$.

$$X^* = (1 + \alpha_2)^{-\frac{1}{2}} [X \quad \sqrt{\alpha_2}I], \quad y^* = [y \quad \mathbf{0}], \quad \gamma = \frac{\alpha_1}{\sqrt{1 + \alpha_2}}, \quad w^* = (\sqrt{1 + \alpha_2})w \quad (1)$$

$$\alpha_1 = \lambda(1 - \alpha), \quad \alpha_2 = \lambda\alpha \quad (2)$$

The elastic net optimization objective function can be written as the following Lasso form[4]:

$$\begin{aligned} L^{Lasso}(\gamma, w) &= \|w^*X^* - y^*\|_2^2 + \gamma\|w^*\|_1 \\ &= \|w^*(1 + \alpha_2)^{-\frac{1}{2}} [X \quad \sqrt{\alpha_2}I] - [y \quad \mathbf{0}]\|_2^2 + \gamma\|w^*\|_1 \\ &= \|w^*(1 + \alpha_2)^{-\frac{1}{2}}X - y\|_2^2 + \frac{\alpha_2}{1 + \alpha_2}\|w^*I\|_2^2 + \gamma\|w^*\|_1 \\ &= \|y\|_2^2 + \frac{\|w^*X\|_2^2}{1 + \alpha_2} - \frac{2w^*X}{\sqrt{1 + \alpha_2}}y^T + \frac{\alpha_2}{1 + \alpha_2}\|w^*I\|_2^2 + \gamma\|w^*\|_1 \\ &= \|y\|_2^2 + \|wX\|_2^2 - 2wXy^T + \alpha_2\|w\|_2^2 + \alpha_1\|w\|_1 \\ &= \|wX - y\|_2^2 + \alpha_2\|w\|_2^2 + \alpha_1\|w\|_1 = L^{Elasticnet}(w) \end{aligned} \quad (3)$$

(ii)

Let $y = (y_1, \dots, y_n)$, X is a $n \times d$ matrix representing the data matrix and $w = (w_1, \dots, w_d)$ representing the weight vector. From the description of the problem, we can know each

$$y_i \sim (x_i^T w, \sigma^2).$$

$$P(w|y, X) \propto P(y|w, X)P(w) = \prod_{i=1}^n \frac{1}{\sqrt{2\pi}\sigma} \exp\left\{-\frac{(y_i - x_i^T w)^2}{2\sigma^2}\right\} \prod_{j=1}^d \frac{1}{\sqrt{2\pi}\tau} \exp\left(-\frac{w_j^2}{2\tau^2}\right) \quad (4)$$

Maximizing the posterior probability is equivalent to maximizing the log posterior probability:

$$\ln P(w|y, X) = -\sum_{i=1}^n \frac{(y_i - x_i^T w)^2}{2\sigma^2} - \sum_{j=1}^d \frac{w_j^2}{2\tau^2} + \text{constant} \quad (5)$$

Maximizing the log posterior probability is equivalent to minimizing the negative of equation (5):

$$\begin{aligned} \min \quad & \sum_{i=1}^n \frac{(y_i - x_i^T w)^2}{2\sigma^2} + \sum_{j=1}^d \frac{w_j^2}{2\tau^2} \\ \iff \min \quad & \|y - X^T w\|_2^2 + \alpha \|w\|_2^2, \alpha = \frac{\sigma^2}{\tau^2} \end{aligned} \quad (6)$$

which is in the form of ridge optimization objective function.

Problem 2

(i)

$$D_{T+1}(i) = \frac{1}{m} \frac{1}{\prod_t Z_t} \exp(-y_i g(x_i)) \quad (7)$$

Proof. Starting from the left hands side:

$$\begin{aligned} D_{T+1}(i) &= \frac{1}{Z_t} D_t(i) \exp(-y_i(a_t f_t(x_i))) \\ &= \frac{1}{Z_t} \frac{1}{Z_{t-1}} D_{t-1}(i) \exp(-y_i(a_{t-1} f_{t-1}(x_i))) \exp(-y_i(a_t f_t(x_i))) \\ &= \frac{1}{Z_t} \frac{1}{Z_{t-1}} D_{t-1}(i) \exp(-y_i(a_{t-1} f_{t-1}(x_i) + a_t f_t(x_i))) \\ &= \frac{1}{Z_t Z_{t-1} Z_{t-2}} D_{t-2}(i) \exp(-y_i(a_{t-2} f_{t-2}(x_i) + a_{t-1} f_{t-1}(x_i) + a_t f_t(x_i))) \end{aligned}$$

We can see that there is an emerging pattern:

$$D_{T+1}(i) = \frac{1}{\prod_t Z_t} D_1(i) \exp(-y_i(\sum_t a_t f_t(x_i)))$$

Recall from the algorithm initialization $D_1(i) = \frac{1}{m}$, therefore:

$$D_{T+1}(i) = \frac{1}{m} \frac{1}{\prod_t Z_t} \exp(-y_i(\sum_t a_t f_t(x_i)))$$

Also we know that $g(x_i) = \sum_t a_t f_t(x_i)$:

$$D_{T+1}(i) = \frac{1}{m} \frac{1}{\prod_t Z_t} \exp(-y_i g(x_i))$$

This gives us the right hand side and concludes the proof. \square

(ii)

Proof. Knowing that:

$$D_{T+1}(i) = \frac{1}{m} \frac{1}{\prod_t Z_t} \exp(-y_i g(x_i))$$

And because the $D_t(i)$ are normalized, all of their sums should be equal to 1, in other words, $\sum_i D_{T+1}(i) = 1$, using this fact we can get the value of $\prod_t Z_t$:

$$\begin{aligned} \sum_i D_{T+1}(i) &= \frac{1}{m} \frac{1}{\prod_t Z_t} \sum_i \exp(-y_i g(x_i)) \\ 1 &= \frac{1}{m} \frac{1}{\prod_t Z_t} \sum_i \exp(-y_i g(x_i)) \\ \prod_t Z_t &= \frac{1}{m} \sum_i \exp(-y_i g(x_i)) \end{aligned}$$

The right hand side is the exponential loss of $g(x)$. Now by using the given fact that the one zero loss is less than the exponential loss:

$$\begin{aligned} \text{err}(g) &= \frac{1}{m} \sum_i \mathbf{1}[y_i \neq \text{sign}(g(x_i))] \\ &\leq \frac{1}{m} \sum_i \exp(-y_i g(x_i)) = \prod_t Z_t \end{aligned}$$

This proves that the error of the aggregate classifier is upper bounded by the the product of Z_t . \square

(iii)

Proof. Knowing that:

$$Z_t = \sum_i D_t(i) \exp(-a_t y_i f_t(x_i)) \quad (8)$$

$$\varepsilon_t = \sum_i D_t(i) \mathbf{1}[y_i \neq f_t(i)] \quad (9)$$

$$1 - \varepsilon_t = \sum_i D_t(i) \mathbf{1}[y_i = f_t(i)] \quad (10)$$

We can get the value of Z_t :

$$\begin{aligned} Z_t &= \sum_i D_t(i) \exp(-a_t y_i f_t(x_i)) \\ &= \sum_{y_i \neq f_t(x_i)} D_t(i) \exp(a_t) + \sum_{y_i = f_t(x_i)} D_t(i) \exp(-a_t) \\ &= \varepsilon_t \exp(a_t) + (1 - \varepsilon_t) \exp(-a_t) \end{aligned}$$

We substitute with the value of $a_t = \frac{1}{2} \ln\left(\frac{1 - \varepsilon_t}{\varepsilon_t}\right)$

$$\begin{aligned}
 Z_t &= \varepsilon_t \exp\left(\frac{1}{2} \ln\left(\frac{1 - \varepsilon_t}{\varepsilon_t}\right)\right) + (1 - \varepsilon_t) \exp\left(-\frac{1}{2} \ln\left(\frac{1 - \varepsilon_t}{\varepsilon_t}\right)\right) \\
 &= \varepsilon_t \sqrt{\frac{1 - \varepsilon_t}{\varepsilon_t}} + (1 - \varepsilon_t) \sqrt{\frac{\varepsilon_t}{1 - \varepsilon_t}} \\
 &= \sqrt{\varepsilon_t^2 \frac{1 - \varepsilon_t}{\varepsilon_t}} + \sqrt{(1 - \varepsilon_t)^2 \frac{\varepsilon_t}{1 - \varepsilon_t}} \\
 &= \sqrt{\varepsilon_t(1 - \varepsilon_t)} + \sqrt{(1 - \varepsilon_t)\varepsilon_t} \\
 &= 2\sqrt{\varepsilon_t(1 - \varepsilon_t)}
 \end{aligned}$$

This gives us the right hand side and concludes the proof. \square

(iv)

Proof. We are given that each iteration t , $\varepsilon_t = \frac{1}{2} - \gamma_t$ Starting with the left hand side:

$$\begin{aligned}
 \prod_t 2\sqrt{\varepsilon_t(1 - \varepsilon_t)} &= \prod_t 2\sqrt{\left(\frac{1}{2} - \gamma_t\right)\left(1 - \left(\frac{1}{2} - \gamma_t\right)\right)} \\
 &= \prod_t 2\sqrt{\left(\frac{1}{2} - \gamma_t\right)\left(\frac{1}{2} + \gamma_t\right)} \\
 &= \prod_t 2\sqrt{\frac{1}{4} - \gamma_t^2} \\
 &= \prod_t \sqrt{4\left(\frac{1}{4} - \gamma_t^2\right)} \\
 &= \prod_t \sqrt{1 - 4\gamma_t^2}
 \end{aligned}$$

Which gives us the left hand side of the inequality, to get the right hand side:

$$\begin{aligned}\prod_t 2\sqrt{\varepsilon_t(1 - \varepsilon_t)} &= \prod_t \sqrt{1 - 4\gamma_t^2} \\ &= \prod_t (1 - 4\gamma_t^2)^{\frac{1}{2}} \\ &\leq \prod_t (\exp(-4\gamma_t^2))^{\frac{1}{2}} \\ &= \prod_t \exp(-2\gamma_t^2) \\ &= \exp(-2 \sum_t \gamma_t^2)\end{aligned}$$

This gives us the right hand side of the inequality and concludes the proof. \square

Problem 3

(i)

With the matrix A , binary vector x_i being hashed to b could be expressed as following:

$$\begin{aligned}
 b &= Ax_i \\
 \begin{bmatrix} b^1 \\ b^2 \\ \dots \\ b^p \end{bmatrix} &= \begin{bmatrix} A^1 \\ A^2 \\ \dots \\ A^p \end{bmatrix} x_i \\
 &= \begin{bmatrix} \left(\sum_{l=1}^n A^{1l} x_i^l \right) \bmod 2 \\ \left(\sum_{l=1}^n A^{2l} x_i^l \right) \bmod 2 \\ \dots \\ \left(\sum_{l=1}^n A^{pl} x_i^l \right) \bmod 2 \end{bmatrix} \quad (11)
 \end{aligned}$$

Take the t -th element of b as an example: $b^t = \left(\sum_{l=1}^n A^{tl} x_i^l \right) \bmod 2$. Since $x_i \in \{0, 1\}^n$, supposing there are n_i nonzero elements in x_i , then $\sum_{l=1}^n A^{tl} x_i^l$ is actually to randomly pick nonzero elements from the n_i nonzero elements in x_i . And $b^t = \left(\sum_{l=1}^n A^{tl} x_i^l \right) \bmod 2$ is in fact describing whether we pick even or odd number of nonzero elements from all the nonzero elements in x_i . Since A^t is generated uniformly at random, the probability of the t -th entry of b , $b^t = 0$ is $\text{Prob}(b^t = 0) = 1/2$ and that of $b^t = 1$ is also $\text{Prob}(b^t = 1) = 1/2$. And this is true for any t between 1 and p . Because different entries of matrix A are independent with each other, $\text{Prob}(x_i \rightarrow b) = \prod_{t=1}^p \text{Prob}(A^t x_i \rightarrow b^t) = \prod_{t=1}^p 1/2 = 1/2^p$.

(ii)

We solve this problem using two methods.

METHOD 1: From part (i), the probability of x_j hashing to any b is $1/2^p$. So the probability of x_j hashing to the particular b that x_i is hashed to is $1/2^p$. In other words, the probability of x_i and x_j hashing to the same vector is $1/2^p$.

METHOD 2: That x_i and x_j hashing to the same vector means $Ax_i = Ax_j$:

$$\begin{bmatrix} (\sum_{l=1}^n A^{1l} x_i^l) \bmod 2 \\ (\sum_{l=1}^n A^{2l} x_i^l) \bmod 2 \\ \dots \\ (\sum_{l=1}^n A^{pl} x_i^l) \bmod 2 \end{bmatrix} = \begin{bmatrix} (\sum_{l=1}^n A^{1l} x_j^l) \bmod 2 \\ (\sum_{l=1}^n A^{2l} x_j^l) \bmod 2 \\ \dots \\ (\sum_{l=1}^n A^{pl} x_j^l) \bmod 2 \end{bmatrix} \quad (12)$$

Then

$$\begin{bmatrix} (\sum_{l=1}^n A^{1l} [x_i^l - x_j^l]) \bmod 2 \\ (\sum_{l=1}^n A^{2l} [x_i^l - x_j^l]) \bmod 2 \\ \dots \\ (\sum_{l=1}^n A^{pl} [x_i^l - x_j^l]) \bmod 2 \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \\ \dots \\ 0 \end{bmatrix} \quad (13)$$

which means $A(x_i - x_j) = \mathbf{0}$.

Supposing x_i and x_j have m different elements. Then any $A^t(x_i - x_j) = 0$ for $(1 \leq t \leq p)$ means to pick even number of elements from the m elements. This is the probability of a random variable X with binomial distribution $\text{Binomial}(m, p_x)$ being even¹. Because each entry of matrix A are generated uniformly at random, the probability of any element in A being 0 is equal to that being 1. That means, for any element in the m elements, the probability of it being picked is equal to that of it not being picked. So $p_x = 1/2$. So:

$$\begin{aligned} \text{Prob}(A^t(x_i - x_j) = 0) &= \text{Prob}(X_{m, p_x} \text{ is even}) \\ &= \frac{1}{2}(1 + (1 - 2p_x)^m) \\ &= \frac{1}{2}(1 + (1 - 2 \times \frac{1}{2})^m) \\ &= \frac{1}{2} \end{aligned} \quad (14)$$

Because different elements of b are independent with each other, $\text{Prob}(A(x_i - x_j) = 0) = \prod_{t=1}^p \text{Prob}(A^t(x_i - x_j) = 0) = \prod_{t=1}^p 1/2 = 1/2^p$.

¹For reference, check <https://math.stackexchange.com/questions/1149270/probability-that-a-random-variable-is-even> and <https://math.stackexchange.com/questions/2541864/hashing-the-cube-binary-matrix-combinatorics>

(iii)

The probability of no collisions among the x_i could be represented as following:

$$\begin{aligned}
 \text{Prob (no collisions)} &= 1 - \text{Prob (exist collisions)} \\
 &\geq 1 - \sum_{1 \leq i < j \leq m} \text{Prob}(x_i, x_j \text{ collide}) \\
 &= 1 - \sum_{1 \leq i < j \leq m} 1/2^p \\
 &= 1 - \binom{m}{2} \frac{1}{2^p} \\
 &= 1 - \frac{m(m-1)}{2} \frac{1}{2^p} \\
 &\geq 1 - \frac{m^2}{2} \frac{1}{2^p}
 \end{aligned} \tag{15}$$

If $p \geq 2 \log_2 m$,

$$\begin{aligned}
 \text{Prob (no collisions)} &\geq 1 - \frac{m^2}{2} \frac{1}{2^p} \\
 &\geq 1 - \frac{m^2}{2} \frac{1}{m^2} \\
 &= 1 - 1/2 \\
 &= 1/2
 \end{aligned} \tag{16}$$

So if $p \geq 2 \log_2 m$, there are no collisions among the x_i with probability at least $1/2$.

Problem 4

iii)

Our final regressor is a neural network that has the following architecture:

1. A fully connected input layer with relu activation of 90 units.
2. Dropout layer with 20% drop out ratio.
3. A hidden fully connected layer with relu activation of 64 units.
4. Dropout layer with 20% drop out ratio.
5. A hidden fully connected layer with relu activation of 32 units.
6. Dropout layer with 20% drop out ratio.
7. Batch normalization layer.
8. Output layer with tanh activation layer bounded by the maximum and minimum year interval of 1 unit.

We ran this model with 50 epochs, batch size of 32 and a 10% validation. And we got 5.4469 training loss and 5.5862 validation loss. For the implementation of this neural network we used Keras[1] with a Tensorflow[2] backend.

The only pre processing we tried is features normalization using Scikit learn[3] StandardScaler, and it improved the accuracy of the neural network.

We reached this model after trying several times with different other models and architectures:

1. The first approach we tried was elastic net with 5 fold cross validation on the penalty regularization hyperparameters, this approach gave us a 6.8 training loss and 6.7 validation loss.
2. We tried ridge regression with 5 fold cross validation on the regularization hyperparameters which also gave us similar results nothing less than 6.6 .
3. And then we normalized the features using Scikit learn StandardScaler which didn't improve the results at all, for both models.
4. And Then we decided to give neural networks a try, with a very simple model with an input layer, 1 hidden relu layer, and a linear activation. which gave us a result around 6 for the validation loss.
5. We kept adding layers: dropout layers to prevent over fitting and a bounded output activation to avoid weird predictions (example, year: 2702?).

References

1. Francois Chollet et al. Keras. <https://keras.io>. 2015.
2. Martin Abadi et al. TensorFlow: Large-Scale Machine Learning on Heterogeneous Systems. Software available from [tensorflow.org](https://www.tensorflow.org/). 2015. url: <https://www.tensorflow.org/>.
3. F. Pedregosa et al. “Scikit-learn: Machine Learning in Python”. In: Journal of Machine Learning Research 12 (2011), pp. 2825–2830.
4. H. Zou and T. Hastie. “Regularization and Variable Selection via the Elastic Net”. In: Journal of the Royal Statistical Society: Series B (Statistical Methodology) 67.2 (2003), pp. 301–320.